

Problem 1.1

Consider the vector $x + c$, when we apply $\text{softmax}(x + c)$, we get,

$$\begin{aligned}\text{softmax}(x_i + c) &= \frac{e^{x_i + c}}{\sum_j e^{x_j + c}} \\ &= \frac{e^{x_i} \cdot e^c}{\sum_j e^{x_j} \cdot e^c} \\ &= \frac{e^{x_i} \cdot e^c}{e^c (\sum_j e^{x_j})} \\ &= \frac{e^{x_i}}{\sum_j e^{x_j}} = \text{softmax}(x_i)\end{aligned}$$

The above proof holds for all $x_i \in x$. Thus, $\text{softmax}(x + c) = \text{softmax}(x) \quad \forall c \in \mathbb{R}$

When we use $c = -\max x_i$ or subtract the maximum value exponent term, we get $e^0 = 1$ and all the other terms will be scaled between 0 and 1, thus avoiding large exponent terms which improves numerical stability and avoids overflows.

Problem 1.2

- The range of each element is between (0,1]. The sum over all elements of $\text{softmax}(x)$ is 1.
- Softmax takes a real-valued vector x and converts it into a probability distribution - where the value of $\text{softmax}(x_i)$ specifies the probability of choosing that value among all other values.
- The first step takes the values and maps it in the positive space between $[0, \infty]$. The second step acts as a normalization step, where we divide the result of the first step with a constant and the third step outputs the probability of the occurrence of the value x_i given all other values.

Problem 1.3

A network without non-linear activation functions will have the x value change according to linear function: $x_{i+1} = W_i x_i + b_i$

When applying to multi-layer neural networks, we have:

$$\begin{aligned}y &= W_n x_n + b_n \\ &= W_n (W_{n-1} x_{n-1} + b_{n-1}) + b_n \\ &= W_n W_{n-1} x_{n-1} + W_n b_{n-1} + b_n \\ &= W' x_{n-1} + b' \\ &= W' (W_{n-2} x_{n-2} + b_{n-2}) + b' \\ &\dots\dots\dots \\ &= Wx + b\end{aligned}$$

which is the same as solving a linear regression problem

Problem 1.4

Consider the sigmoid activation function, we get,

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{1 + e^x}$$

Taking the derivative of the function, we get,

$$\begin{aligned}\frac{d\sigma(x)}{dx} &= \frac{d}{dx} \left(\frac{e^x}{1 + e^x} \right) \\&= \frac{e^x(1 + e^x) - e^x e^x}{(1 + e^x)^2} \\&= \frac{e^x}{(1 + e^x)^2} \\&= \frac{e^x}{(1 + e^x)} \cdot \frac{e^x}{(1 + e^x)} \\&= \frac{e^x}{(1 + e^x)} \cdot \left(\frac{1 + e^x - e^x}{(1 + e^x)} \right) \\&= \frac{e^x}{(1 + e^x)} \cdot \left(1 - \frac{e^x}{(1 + e^x)} \right) \\&= \sigma(x) (1 - \sigma(x))\end{aligned}$$

Problem 1.5

$$\frac{dJ}{dy} = \delta$$

$$\frac{dJ}{dy_j} = \delta \quad \text{for } j \text{ in } 1 \text{ to } k$$

$$\frac{dJ}{dW_{ij}} = \frac{dJ}{dy_j} \frac{dy_j}{dW_{ij}} = \delta_j x_i$$

$$\frac{dJ}{dx_i} = \frac{dJ}{dy_j} \frac{dy_j}{dx_i} = \sum W_{ij} * \delta_j$$

$$\frac{dJ}{db_j} = \frac{dJ}{dy_j} \frac{dy_j}{db_j} = \delta_j$$

Thus, forming the matrices, we get,

$$\frac{dJ}{dW} = x\delta^T$$

$$\frac{dJ}{dx} = W\delta$$

$$\frac{dJ}{db} = \delta$$

Problem 1.6

- When we take the sigmoid function, it maps the input space which is in \mathbb{R} to $[0,1]$. As a result of this, there are large regions of the input space which are mapped to an extremely small range. In these regions of the input space, even a large change in the input will produce a small change in the output - hence the gradient is small. Consider the gradient of the sigmoid between $[0,1]$, it is constrained between $[0,0.25]$
- When the gradient is propagated through the network, the gradients become smaller and smaller, thus after several layers, the gradients almost “vanish”, that is they become zero quickly - resulting in almost no change in the parameters of the network, making it difficult to train the network.
- The output range of the sigmoid function is $[0,1]$ and the output range of the tanh function is $[-1,1]$. Choosing the tanh function is more preferable as it maps the input space which is in \mathbb{R} to a large range of values and also, it maps positive values to positive values and negative values to negative values.
- The derivative values of $\tanh(x)$ are constrained between $[0,1]$ for values between $[-1,1]$, thus having stronger and larger derivatives for the same inputs compared to the sigmoid function. Thus, $\tanh(x)$ would have lesser of a vanishing gradient problem compared to $\text{sigmoid}(x)$.
- To make $\tanh(x)$ have the range of $\text{sigmoid}(x)$,

$$\tanh(x) = \frac{1 - e^{-2x}}{1 + e^{-2x}}$$

$$\tanh(x/2) = \frac{1 - e^{-x}}{1 + e^{-x}}$$

$$\tanh(x/2) + 1 = \frac{1 - e^{-x}}{1 + e^{-x}} + 1$$

$$= \frac{2}{1 + e^{-x}}$$

$$= 2\sigma(x)$$

So, $\tanh(x) + 1 = 2\sigma(2x)$

And $\tanh(x) = 2\sigma(2x) - 1$

Problem 2.1.1

Initializing all the weights to zero does not prove to be useful. If we do so, the weights we learn will be the same for the same values of inputs across iterations and thus, the network will not be able to effectively map the input space to the output space.

In the event we take a network with a few linear activations (especially in the last few layers) and also take the weights and biases to be zero, then the network will essentially not learn anything and the output of the network would always be zero.

Problem 2.1.2

```
6 ##### Q 2.1 #####
7 # initialize b to 0 vector
8 # b should be a 1D array, not a 2D array with a singleton dimension
9 # we will do  $XW + b$ .
10 # X be [Examples, Dimensions]
11 def initialize_weights(in_size, out_size, params, name=''):
12
13     bound = np.sqrt(6) / np.sqrt(in_size + out_size)
14     W, b = np.random.uniform(-bound, bound, (in_size, out_size)), np.zeros(out_size)
15
16     params['W' + name] = W
17     params['b' + name] = b
18
```

Problem 2.1.3

Random weights help as they allow for the network to compute different updates for different units and this ensures that no input patterns are lost in null space of the forward propagation and no gradient patterns are lost in the null space of the backward propagation.

Scaling the initialization based on the layer size helps to keep the information flowing in the network, by allowing the variance of the activations in subsequent layers be the same during forward propagation and the variance of the weights in subsequent layers be same during backward propagation.

Problem 2.2.1

```
19 ##### Q 2.2.1 #####
20 # x is a matrix
21 # a sigmoid activation function
22 def sigmoid(x):
23
24     res = 1/(1+np.exp(-x))
25
26     return res
27
28 ##### Q 2.2.1 #####
29 def forward(X,params,name='',activation=sigmoid):
30     """
31     Do a forward pass
32
33     Keyword arguments:
34     X -- input vector [Examples x D]
35     params -- a dictionary containing parameters
36     name -- name of the layer
37     activation -- the activation function (default is sigmoid)
38     """
39     pre_act, post_act = None, None
40     # get the layer parameters
41     W = params['W' + name]
42     b = params['b' + name]
43
44
45     pre_act = np.matmul(X, W) + b
46     post_act = activation(pre_act)
47
48
49     # store the pre-activation and post-activation values
50     # these will be important in backprop
51     params['cache_' + name] = (X, pre_act, post_act)
52
53     return post_act
```

Problem 2.2.2

```
##### Q 2.2.2 #####
# x is [examples,classes]
# softmax should be done for each row
def softmax(x):

    ss = np.max(x, axis=1)
    ss = ss[:, np.newaxis]
    ex = np.exp(x-ss)
    div = np.sum(ex, axis=1)
    div = div[:, np.newaxis]
    res = ex/div

    return res
```

Problem 2.2.3

```
##### Q 2.2.3 #####
# compute total loss and accuracy
# y is size [examples,classes]
# probs is size [examples,classes]
def compute_loss_and_acc(y, probs):
    loss, acc = None, 0

    loss = -np.sum(y*np.log(probs))

    for i in range(y.shape[0]):
        if(np.argmax(y[i, :]) == np.argmax(probs[i, :])):
            acc+=1
    acc = acc/y.shape[0]

    return loss, acc
```

Problem 2.3

```
94 def backwards(delta,params,name='',activation_deriv=sigmoid_deriv):
95     """
96     Do a backwards pass
97
98     Keyword arguments:
99     delta -- errors to backprop
100     params -- a dictionary containing parameters
101     name -- name of the layer
102     activation_deriv -- the derivative of the activation_func
103     """
104     grad_X, grad_W, grad_b = None, None, None
105     # everything you may need for this layer
106     W = params['W' + name]
107     b = params['b' + name]
108     X, pre_act, post_act = params['cache_' + name]
109
110     # do the derivative through activation first
111     # (don't forget activation_deriv is a function of post_act)
112     # then compute the derivative W, b, and X
113
114     grad_op = activation_deriv(post_act)*delta
115
116     grad_W = np.matmul(X.T, grad_op)
117     grad_b = np.sum(grad_op, axis=0)
118     grad_X = np.matmul(W, grad_op.T).T
119
120
121     # store the gradients
122     params['grad_W' + name] = grad_W
123     params['grad_b' + name] = grad_b
124     return grad_X
125
```

Problem 2.4

```
126 ##### Q 2.4 #####
127 # split x and y into random batches
128 # return a list of [(batch1_x, batch1_y)...]
129 def get_random_batches(x, y, batch_size):
130     batches = []
131
132     random_batches = np.array_split(np.random.permutation(x.shape[0]), x.shape[0]//batch_size)
133
134     for i in random_batches:
135         batches.append((x[i, :], y[i, :]))
136
137     return batches
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182 # WRITE A TRAINING LOOP HERE
183 max_iters = 500
184 learning_rate = 1e-3
185 # with default settings, you should get loss < 35 and accuracy > 75%
186 for itr in range(max_iters):
187     total_acc = 0
188     total_loss = 0
189     avg_acc = 0
190     for xb, yb in batches:
191         #####
192         ##### your code here #####
193         #####
194         pass
195         # forward
196         h1 = forward(xb, params, 'layer1', sigmoid)
197         probs = forward(h1, params, 'output', softmax)
198
199         # loss
200         loss, acc = compute_loss_and_acc(yb, probs)
201
202         # be sure to add loss and accuracy to epoch totals
203         total_loss += loss
204         total_acc += acc
205
206         # backward
207         delta = probs - yb
208         delta = backwards(delta, params, 'output', linear_deriv)
209         backwards(delta, params, 'layer1', sigmoid_deriv)
210
211         # apply gradient
212         # gradients should be summed over batch samples
213         for layer in ['output', 'layer1']:
214             params['W' + layer] -= learning_rate * params['grad_W' + layer]
215             params['b' + layer] -= learning_rate * params['grad_b' + layer]
216
217     avg_acc = total_acc / batch_num
218
219     if itr % 100 == 0:
220         print("itr: {:02d} \t loss: {:.2f} \t acc : {:.2f}".format(itr, total_loss, avg_acc))
221
```




Problem 2.5

```
138 # compute gradients using finite difference
139 eps = 1e-6
140 for k,v in params.items():
141     if '_' in k:
142         continue
143
144     if 'W' in k:
145         for r in range(v.shape[0]):
146             for c in range(v.shape[1]):
147                 d = np.zeros(v.shape)
148                 d[r, c] = 1
149
150                 params1 = copy.deepcopy(params)
151                 params1[k] = v + eps*d
152                 h1 = forward(xb, params1, 'layer1', sigmoid)
153                 probs = forward(h1, params1, 'output', softmax)
154                 f1, _ = compute_loss_and_acc(yb, probs)
155
156                 params2 = copy.deepcopy(params)
157                 params2[k] = v - eps * d
158                 h1 = forward(xb, params2, 'layer1', sigmoid)
159                 probs = forward(h1, params2, 'output', softmax)
160                 f2, _ = compute_loss_and_acc(yb, probs)
161
162                 params['grad_' + k][r, c] = (f1 - f2) / (2*eps)
163
164     # Bias
165     else:
166         for i in range(v.shape[0]):
167             d = np.zeros(v.shape)
168             d[i] = 1
169             params1 = copy.deepcopy(params)
170             params1[k] = v + eps*d
171             h1 = forward(xb, params1, 'layer1', sigmoid)
172             probs = forward(h1, params1, 'output', softmax)
173             f1, _ = compute_loss_and_acc(yb, probs)
174
175             params2 = copy.deepcopy(params)
176             params2[k] = v - eps*d
177             h1 = forward(xb, params2, 'layer1', sigmoid)
178             probs = forward(h1, params2, 'output', softmax)
179             f2, _ = compute_loss_and_acc(yb, probs)
180
181             params['grad_' + k][i] = (f1 - f2) / (2*eps)
182
183 total_error = 0
184 for k in params.keys():
```

Problem 3.1

```
PS C:\Users\mathe\OneDrive\Desktop\CV\hw5\hw5\python> python .\run_q3.py
itr: 00      loss: 34852.33      acc : 0.18
itr: 02      loss: 18811.02      acc : 0.57
itr: 04      loss: 14243.28      acc : 0.65
itr: 06      loss: 12370.75      acc : 0.69
itr: 08      loss: 11235.84      acc : 0.72
itr: 10      loss: 10406.66      acc : 0.74
itr: 12      loss: 9737.54      acc : 0.75
itr: 14      loss: 9165.39      acc : 0.77
itr: 16      loss: 8658.59      acc : 0.78
itr: 18      loss: 8199.68      acc : 0.79
itr: 20      loss: 7778.33      acc : 0.80
itr: 22      loss: 7388.05      acc : 0.81
itr: 24      loss: 7024.56      acc : 0.82
itr: 26      loss: 6684.85      acc : 0.83
itr: 28      loss: 6366.69      acc : 0.84
itr: 30      loss: 6068.27      acc : 0.85
itr: 32      loss: 5787.98      acc : 0.86
itr: 34      loss: 5524.35      acc : 0.87
itr: 36      loss: 5276.06      acc : 0.87
itr: 38      loss: 5041.87      acc : 0.88
itr: 40      loss: 4820.70      acc : 0.88
itr: 42      loss: 4611.56      acc : 0.89
itr: 44      loss: 4413.60      acc : 0.90
itr: 46      loss: 4226.04      acc : 0.90
itr: 48      loss: 4048.20      acc : 0.91
Validation accuracy: 0.7916666666666666
Test accuracy: 0.7905555555555555
```



The validation accuracy of 79% was obtained with **learning rate = 0.003 or 3e-3** and **batch size=5**. The plots of loss and accuracy of the model trained for 50 epochs are shown below.



Problem 3.2



1) Highest Accuracy

The validation accuracy of 79% was obtained with **learning rate = 0.003 or 3e-3** and **batch size=5**. The plots of loss and accuracy of the model trained for 50 epochs are shown below.




2) Higher Learning Rate

When the learning rate was 0.03 (10 times higher), the losses were higher than the previous case while the accuracy was lower. The curves were also bumpier because the step size was too large.




3) Lower Learning Rate


When the learning rate was 0.0003 (1/10 times), the curves are smooth as those shown, but because the step size was too small, it did not converge to the optimum within the maximum number of iterations (50) that we set.



Problem 3.3


Layer 1 weights after initialization





On comparing the two figures, we can see that the initial weights have no patterns in them as we initialized the layers with random uniform distribution. The weights after 50 epoch training have some more clear patterns, visible strokes of the letters and numbers which look somewhat averaged over the dataset

Problem 3.4



From the above confusion matrix, it can be seen that the brighter the grid, the greater the number of correctly predicted labels in the grid. Since the main diagonal is the heaviest, we can conclude that the results are pretty good.

After training 50 epochs the top few pairs that are commonly confused are number '0' and letter 'O', '2' and 'Z', '5' and 'S', '4' and 'Y', '0' and 'D', which are often misjudged manually too.

Problem 4.1

The assumptions the sample method makes are as follows:

- 1) The stroke width, style of writing of the characters extracted are assumed to be similar to the NIST36 dataset that were trained. Also, the pixel values of the extracted characters are assumed to be similar to the ones in the dataset, however the dataset does not have perfect binary black-and-white images upon inspection.

The image shows three handwritten characters: 'H', 'T', and 'E'. The 'H' has a thick vertical stroke on the left and a thinner one on the right. The 'T' has a thick vertical stem and a thin horizontal top bar. The 'E' has a thick vertical stem and thin horizontal bars. This illustrates a variation in stroke width from the NIST36 dataset.

From the above image, the stroke widths are different than in the dataset and these images might fail to be detected correctly

- 2) The cropped image is assumed to not have any noises, any overlap between two letters and parts of letters to be fully connected.

The image shows two handwritten characters: 'R' and 'IT'. The 'R' is written with a single stroke. The 'IT' is written with the 'I' and 'T' overlapping, where the vertical stroke of the 'I' is partially covered by the vertical stroke of the 'T'. This illustrates a failure of the assumption that letters are fully connected without overlap.

In the above image, the first letter has incomplete parts, and second letter and third letters are overlapping. This image might fail to be detected correctly

Problem 4.2



```
# find the rows using..RANSAC, counting, clustering, etc.
heights = [bbox[2]-bbox[0] for bbox in bboxes]
mean_height = sum(heights)/len(heights)



centers = [((bbox[2]+bbox[0])/2, (bbox[3]+bbox[1])/2, bbox[2]-bbox[0], bbox[3]-bbox[1]) for bbox in bboxes]
centers = sorted(centers, key=lambda p: p[0])
rows = []
pre_h = centers[0][0]

row = []
for c in centers:
    if c[0] > pre_h + mean_height:
        row = sorted(row, key = lambda p:p[1])
        rows.append(row)
        row = [c]
        pre_h = c[0]
    else:
        row.append(c)
row = sorted(row, key = lambda p:p[1])
rows.append(row)

# crop the bounding boxes
# note.. before you flatten, transpose the image (that's how the dataset is!)
# consider doing a square crop, and even using np.pad() to get your images looking more like the dataset
kernel = np.array([[0, 1, 0], [1, 1, 1], [0, 1, 0]])
data = []
for row in rows:
    row_data = []
    for y, x, h, w in row:
        # crop out the character
        crop = bw[y-h//2:y+h//2, x-w//2:x+w//2]
        # pad it to square
        h_pad, w_pad = 0, 0
        if h > w:
            h_pad = h//20
            w_pad = (h-w)//2+h_pad
        elif h < w:
            w_pad = w//20
            h_pad = (w-h)//2+w_pad
        crop = np.pad(crop, ((h_pad, h_pad), (w_pad, w_pad)), 'constant', constant_values=(1, 1))
        # resize to 32*32
        crop = skimage.transform.resize(crop, (32, 32))
        crop = skimage.morphology.erosion(crop, kernel)
        crop = np.transpose(crop)
        row_data.append(crop.flatten())
    data.append(np.array(row_data))
```


Problem 4.3





From the above images it can be seen that the algorithm was able to detect all the letters in the given images with 100% accuracy.

Problem 4.4

The extracted text and its accuracy from the detection is as below:

```
TQDQLIST  
IMAKEATDQDLIST  
2LHFCKDFF7HEFIRFWT  
THINGQNTQDQLIST  
3RIALIZEY0UHAUEALR6ADT  
CQMPFLT5DITHINGS  
9RWARDYDU8GELFWITH  
ANAP
```

```
ABCDEFGH  
IJKLMN  
OPQRSTW  
VWXYZ  
123GS6789J
```

```
HAIKUSARHHAGY  
BLTSDMETIMESTHEYDDWTMAKGBHNGE  
RBGRIGERAMQR
```

```
JEEPLKAKMING  
DEPPEKLEARNING  
DEBPE5TLEARNING
```

We can see that, overall, the results were pretty good and the accuracy is pretty decent as well.

Problem 5.1.1 and 5.1.2

```
# Q5.1 & Q5.2
# initialize layers here
initialize_weights(train_x.shape[1],hidden_size,params,'layer1')
initialize_weights(hidden_size,hidden_size,params,'layer2')
initialize_weights(hidden_size,hidden_size,params,'layer3')
initialize_weights(hidden_size,train_x.shape[1],params,'output')

keys = [key for key in params.keys()]
for k in keys:
    params['m_'+k] = np.zeros(params[k].shape)
train_loss=[]

# should look like your previous training loops
losses = []
for itr in range(max_iters):
    total_loss = 0
    for xb,_ in batches:
        # training loop can be exactly the same as q2!
        # your loss is now squared error
        # delta is the d/dx of (x-y)^2
        # to implement momentum
        # just use 'm_'+name variables
        # to keep a saved value over timestamps
        # params is a Counter(), which returns a 0 if an element is missing
        # so you should be able to write your loop without any special conditions


        h1 = forward(xb, params, 'layer1', relu)
        h2 = forward(h1, params, 'layer2', relu)
        h3 = forward(h2, params, 'layer3', relu)
        probs = forward(h3, params, 'output', sigmoid)

        # Loss and accuracy
        loss = np.sum((xb - probs)**2)
        total_loss += loss

        # Backward pass
        delta = 2*(probs-xb)
        delta = backwards(delta, params, 'output', sigmoid_deriv)
        delta = backwards(delta, params, 'layer3', relu_deriv)
        delta = backwards(delta, params, 'layer2', relu_deriv)
        backwards(delta, params, 'layer1', relu_deriv)

        # Apply gradient
        for layer in ['output','layer1','layer2','layer3']:
            params['m_W' + layer] = 0.9*params['m_W' + layer] - learning_rate * params['grad_W' + layer]
            params['W' + layer] += params['m_W' + layer]
            params['m_b' + layer] = 0.9*params['m_b' + layer] - learning_rate * params['grad_b' + layer]
            params['b' + layer] += params['m_b' + layer]
```

Problem 5.2



We can see that initially there is a drastic drop in the total loss of the system, but as the number of epochs increase, the momentum of convergence decreases and the total loss starts decreases slowly. Also, since we are changing the learning rate as iterations increase, we also slow down the convergence of the system, if the learning rate becomes relatively small.


Problem 5.3.1



Problem 5.3.2


The average PSNR over the validation data set is **14.581606637265628**

Problem 6.1.1




Validation accuracy: 0.7741666666666667

Problem 6.1.2



Validation accuracy: 0.9027777777777778

Problem 6.1.3



Problem 6.2

```
Starting epoch 22 / 24
Train accuracy: 0.8661764705882353
Val accuracy: 0.788235294117647
Starting epoch 23 / 24
Train accuracy: 0.8544117647058823
Val accuracy: 0.7617647058823529
Starting epoch 24 / 24
Train accuracy: 0.8838235294117647
Val accuracy: 0.8117647058823529
```

```
Train Epoch: 21, Loss: 2.314396
Train Epoch: 21, Loss: 2.284284
Train Epoch: 21, Loss: 2.378585
Test set: Average loss: 2.1330, Accuracy: 25.88%
Train Epoch: 22, Loss: 2.102735
Train Epoch: 22, Loss: 2.061395
Train Epoch: 22, Loss: 2.200077
Test set: Average loss: 2.1505, Accuracy: 25.88%
Train Epoch: 23, Loss: 2.257583
Train Epoch: 23, Loss: 2.289290
Train Epoch: 23, Loss: 1.847931
Test set: Average loss: 2.2278, Accuracy: 20.59%
```

The fine-tuned network and the scratch network are implemented in q7-finetune.py. The maximum validation accuracy for the finetuned network is **81.176%** and the maximum validation accuracy for the scratch trained network is **25.88%**. We can see that the finetuned network outperforms the scratch trained network.

Since the model is using pre-trained weights (from ImageNet dataset = very large number of images for it to learn from), it has better understanding (more semantic information of how images in scenes and flowers in such scenes can look like and can build better representations internally for the image - making it easier for it to identify and classify the image).